

Fitting copulas

The goal of this project is to study a concrete example of multivariate stock returns using copula estimation. More specifically, we will:

- Fit univariate models for each margin individually.
- Fit a copula to account for the dependence between the margins.
- Simulate new data following the fitted model and use it to make predictions about the aggregated loss.

0) Setup

The first step is to load the necessary libraries.

```
library(rugarch)
library(xts)
library(ADGofTest)
library(qqtest)
library(copula)
library(qrmdata)
library(qrmtools)

set.seed(123) # for reproducibility
```

1) Data Preparation:

The data consists of an `xts` object containing adjusted close prices of the constituents of the S&P 500 index. The data has been obtained from Yahoo Finance as of 2016-01-03, via the function `get_data()` from the package `qrmtools`. The constituents data ranges from the first date at least one of the constituents is available (with missing data if not available) to 2015-12-31.

We will only consider 5 constituents of the index: Intel, Qualcomm, Google, Apple and Microsoft from January 3rd, 2007 to December 31st, 2009.

First, we create a matrix `S` containing the observations of interest and check that it doesn't contain missing data.

```
# load the constituents data of the S&P 500
data("SP500_const")
stocks <- c("INTC", "QCOM", "GOOGL", "AAPL", "MSFT")
time <- c("2007-01-03", "2009-12-31")

#observations of interest:
S <- SP500_const[paste0(time, collapse = "/"), stocks]

## Check for missing data
any(is.na(S))
```

```
## [1] FALSE
```

Then, we compute the negative log-returns (representing financial losses)

```
## Build -log-returns
X <- -returns(S,method = 'logarithm')

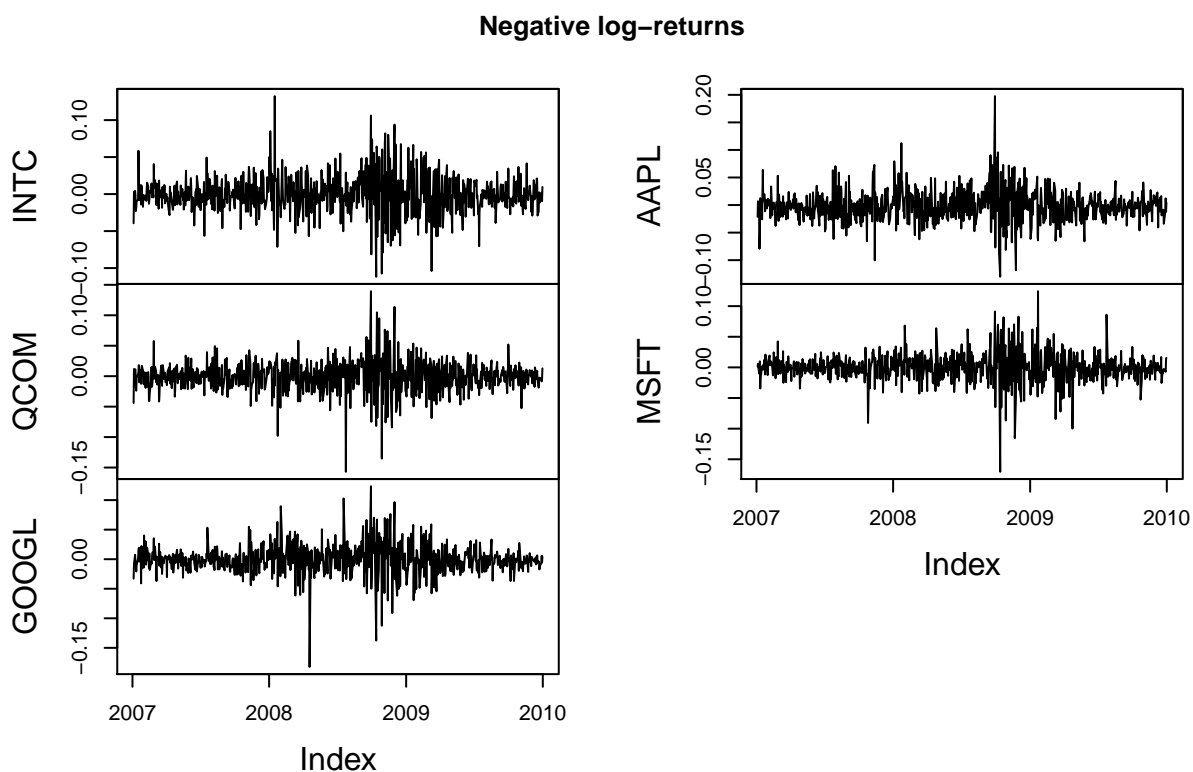
# sample size
n <- nrow(X)

# dimension
p <- ncol(X)
```

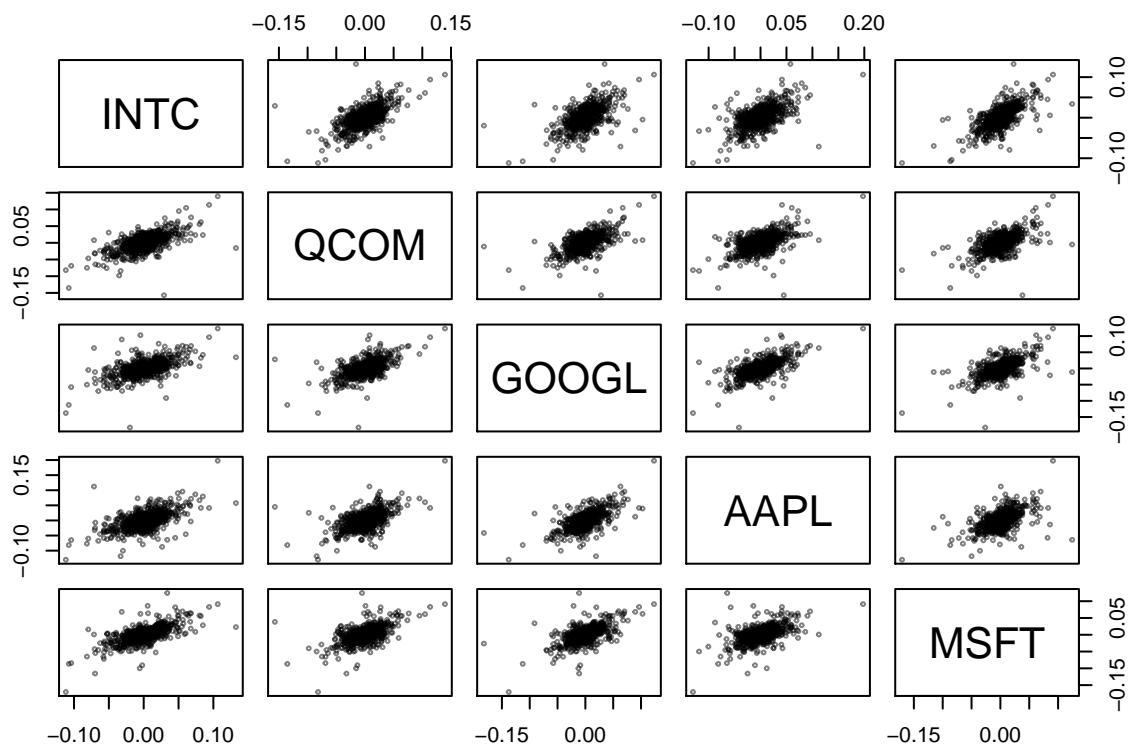
We now create some exploratory plots to study their evolution in time (acf of X and X^2), their marginal distributions, their joint behavior.

```
## Basic plot: use functions plot.zoo, pairs, histograms
# eventually draw acf, pacf for each margin

plot.zoo(X, main = "Negative log-returns")
```



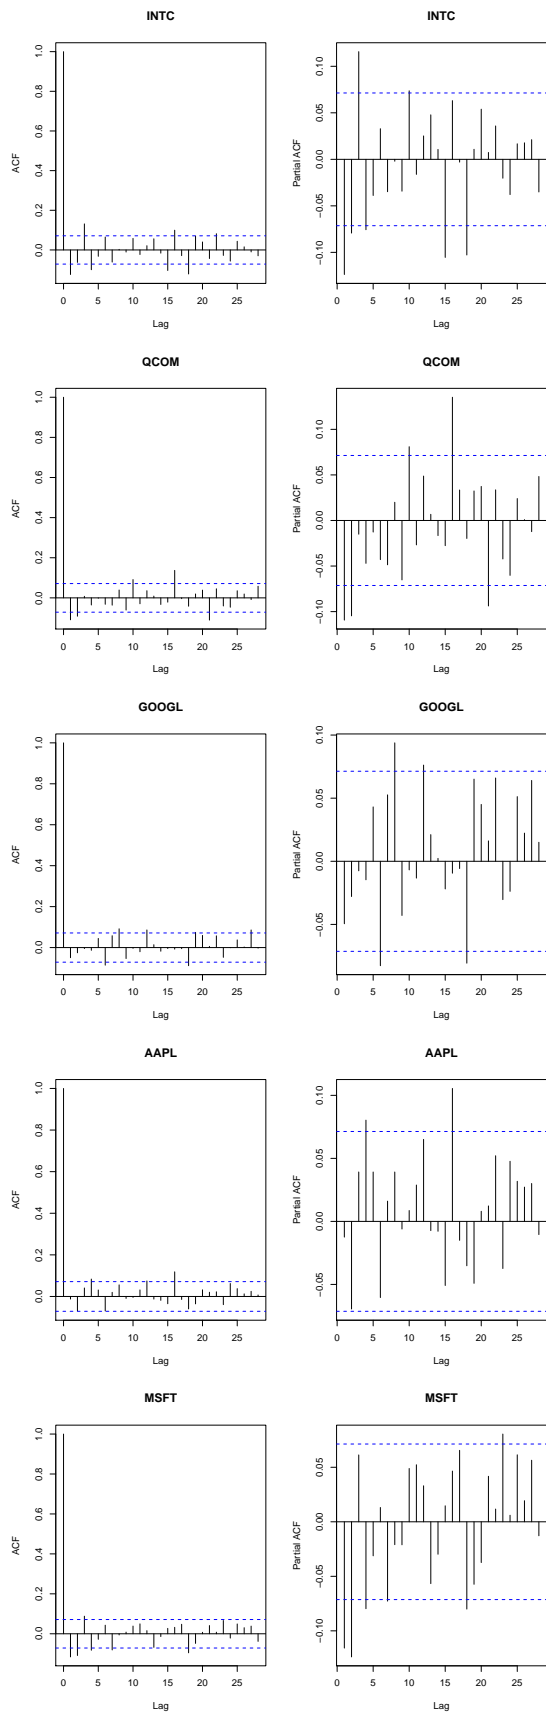
```
pairs(as.matrix(X),cex = 0.4, col = adjustcolor("black", alpha.f = 0.5))
```



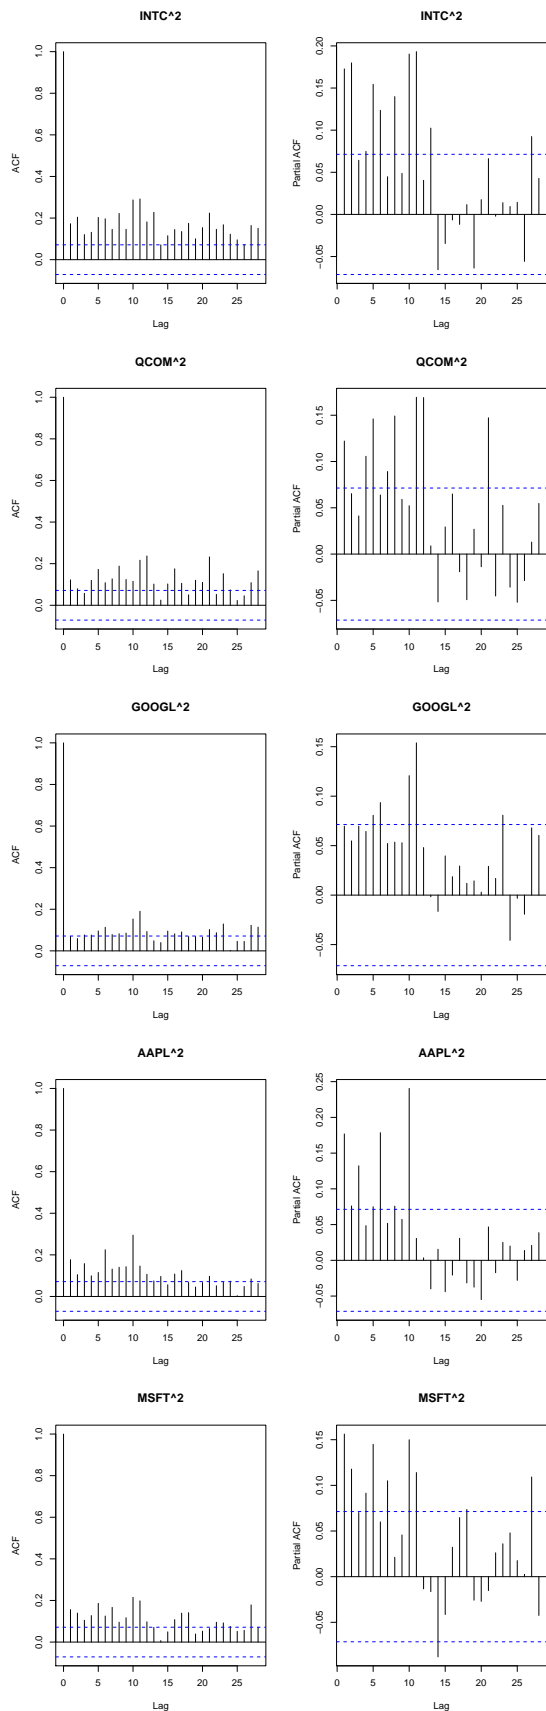
```

par(mfrow=c(5,2))
for (k in 1:5)
{
  acf(X[,k],main=colnames(X)[k])
  pacf(X[,k],main=colnames(X)[k])
}

```



```
par(mfrow=c(5,2))
for (k in 1:5)
{
  acf(X[,k]^2,main=paste0(colnames(X)[k],'^2'))
  pacf(X[,k]^2,main=paste0(colnames(X)[k],'^2'))
}
```



We want to model the dependence in the mean and the variance of the process $(X_t)_{t>0}$.

The autocorrelation function of $(X_t)_{t>0}$ shows some evidence of serial correlation at small lags. This suggests that an ARMA(1,1) could be suitable for modelling.

The autocorrelation functions of the squared components $(X_t^2)_{t>0}$, representing the spot volatility (under the assumption of zero mean) also shows us that the volatility is heteroskedastic. We can take this into account through a GARCH(1,1) process.

In fact, such a model is often used for time series of financial returns.

2) Fitting marginal models

We consider an ARMA(1,1)-GARCH(1,1) with standardized t residuals. This model is fitted by using the function `fit_ARMA_GARCH` in the package **qrmtools**.

```
## Fit marginal time series
uspec <- rep(list(ugarchspec(distribution.model = "std")), ncol(X))
fit.ARMA.GARCH <- fit_ARMA_GARCH(X, ugarchspec.list = uspec)
```

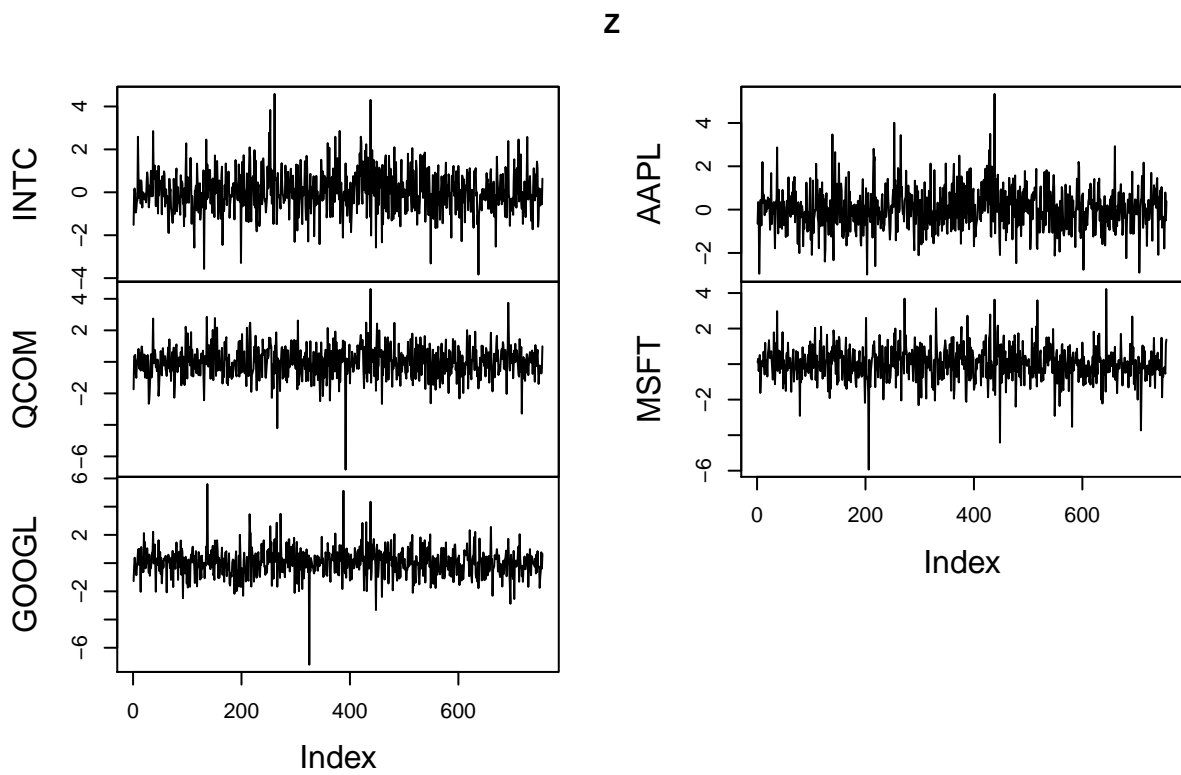
```
## =====
```

```
# fitted models
fits <- fit.ARMA.GARCH$fit
```

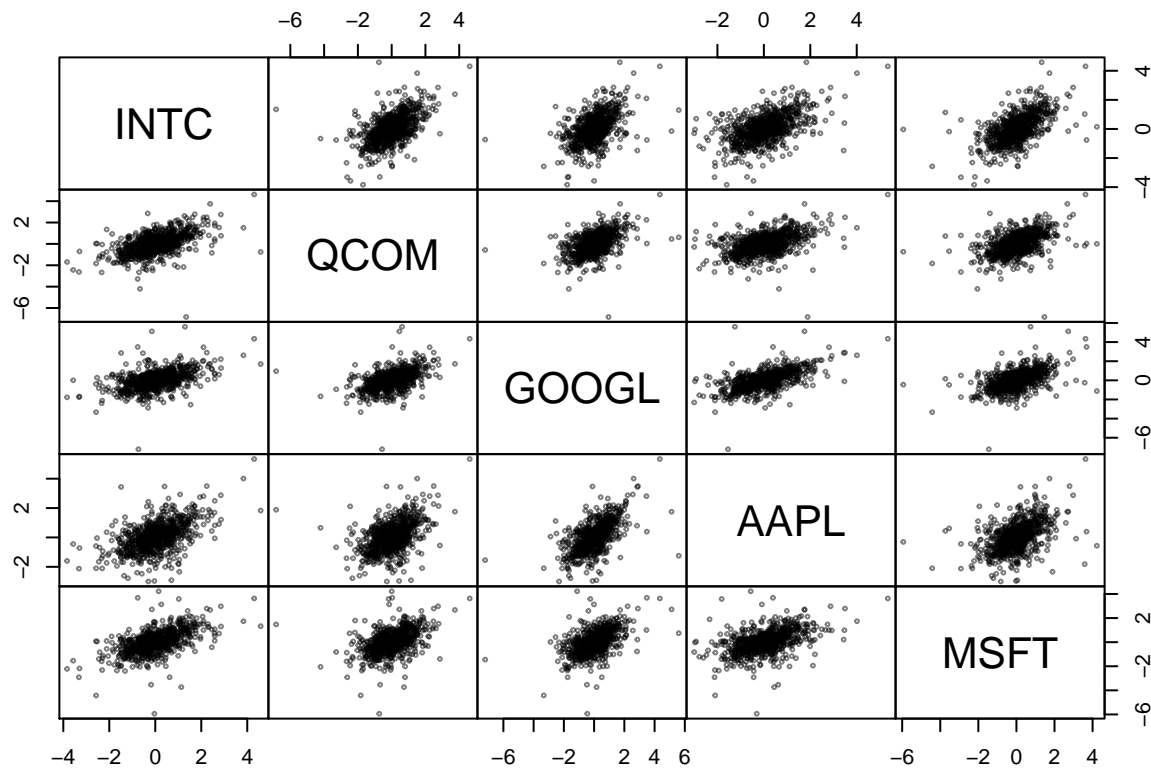
Next we compute the standardized residuals for each margin, and plot these

```
# a matrix n x p containing the standardized marginal residuals
Z <- as.matrix(do.call(merge, lapply(fits, residuals, standardize = TRUE))) # grab out standardized residuals
colnames(Z) <- colnames(S)

# visualize the standardized residuals
plot.zoo(Z)
```



```
pairs2(Z,cex = 0.4, col = adjustcolor("black", alpha.f = 0.5))
```

And then extract the estimated degrees of freedom for each margin $\hat{\nu}_1, \dots, \hat{\nu}_5$.

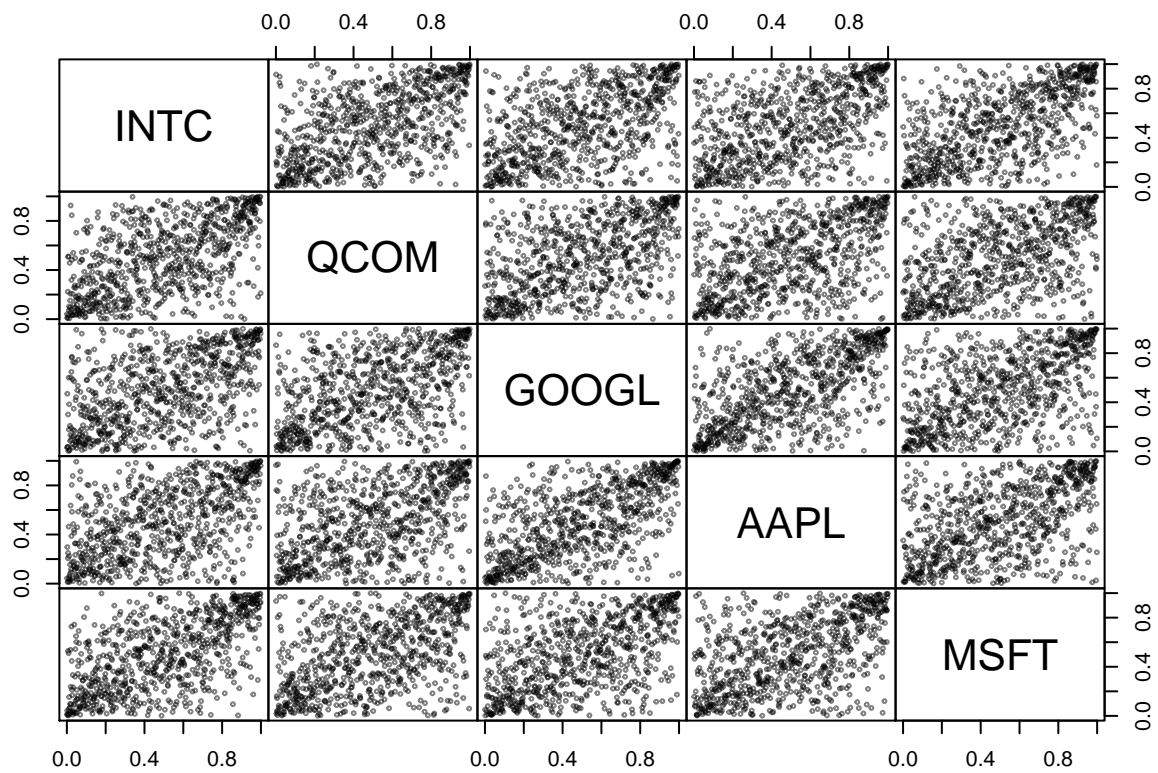
```
# vector of estimated df
(nu.mar <- vapply(fits, function(x) x@fit$coef[["shape"]], NA_real_))

## [1] 6.093672 7.151326 4.822795 7.508152 4.532731
```


3) Fitting copulas

Pseudo-observations from the standardized t residuals are then generated, and plotted, through two methods; non-parametrically (based on the ranks) or parametrically (based on the estimated marginal distribution of the residuals, $t_{\hat{\nu}_i}$).

```
## Compute pseudo-observations from the standardized t residuals
U <- pobs(Z)
pairs2(U, cex = 0.4, col = adjustcolor("black", alpha.f = 0.5))
```



```
# alternatively
U2<- pt(Z,nu.mar)
pairs2(U2, cex = 0.4, col = adjustcolor("black", alpha.f = 0.5))
```



Now, a series of different copulas will be fitted through the use of the `fitCopula` function from the **copula** package.

Starting with a Gumbel copula on the pseudo-observations.

```
## Fitting a Gumbel copula
fit.gc <- fitCopula(gumbelCopula(dim = p),
  data = U, method = "mpl")
```

```
# print estimated copula parameter
fit.gc@estimate
```

```
## [1] 1.454779
```

```
# fitted copula
gc <- fit.gc@copula
```

The matrices of pairwise Kendall's tau and upper tail-dependence coefficients (which are both the same for all non-diagonal elements, which follows from the fact that the Gumbel copula is archimedean) based on the fitted copula are then the following:

```
## Compute matrices of pairwise Kendall's tau and upper tail-dependence coefficients
p2P(tau(gc), d = p)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 1.0000000 0.3126104 0.3126104 0.3126104 0.3126104
## [2,] 0.3126104 1.0000000 0.3126104 0.3126104 0.3126104
## [3,] 0.3126104 0.3126104 1.0000000 0.3126104 0.3126104
## [4,] 0.3126104 0.3126104 0.3126104 1.0000000 0.3126104
```

```
## [5,] 0.3126104 0.3126104 0.3126104 0.3126104 1.0000000
```

```
p2P(lambda(gc)["upper"], d = p)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 1.0000000 0.3896329 0.3896329 0.3896329 0.3896329
## [2,] 0.3896329 1.0000000 0.3896329 0.3896329 0.3896329
## [3,] 0.3896329 0.3896329 1.0000000 0.3896329 0.3896329
## [4,] 0.3896329 0.3896329 0.3896329 1.0000000 0.3896329
## [5,] 0.3896329 0.3896329 0.3896329 0.3896329 1.0000000
```

Next, the same procedure follows for a t copula. Note that the tau and tail dependence coefficients are no longer all the same this time.

```
## Fitting a t copula
fit.tc <- fitCopula(tCopula(dim = p, dispstr = "un"),
                    data = U, method = "itau.mpl")
# estimated degrees of freedom nu
(nu <- tail(fit.tc@estimate, n = 1))
```

```
## [1] 7.269678
```

```
# estimated correlation matrix
(P <- p2P(head(fit.tc@estimate, n = -1)))
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 1.0000000 0.5761937 0.5246273 0.5361455 0.6246738
## [2,] 0.5761937 1.0000000 0.5229914 0.5064793 0.5123869
## [3,] 0.5246273 0.5229914 1.0000000 0.6429929 0.5391979
## [4,] 0.5361455 0.5064793 0.6429929 1.0000000 0.5249092
## [5,] 0.6246738 0.5123869 0.5391979 0.5249092 1.0000000
```

```
# fitted copula
tc <- fit.tc@copula
```

```
## Compute matrices of pairwise Kendall's tau and upper tail-dependence coefficients
p2P(tau(tc))
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 1.0000000 0.3909252 0.3515906 0.3602403 0.4295361
## [2,] 0.3909252 1.0000000 0.3503680 0.3381067 0.3424772
## [3,] 0.3515906 0.3503680 1.0000000 0.4446150 0.3625450
## [4,] 0.3602403 0.3381067 0.4446150 1.0000000 0.3518014
## [5,] 0.4295361 0.3424772 0.3625450 0.3518014 1.0000000
```

```
p2P(lambda(tc)[(choose(p,2)+1):(p*(p-1))])
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 1.0000000 0.1730450 0.1457634 0.1514811 0.2030989
## [2,] 0.1730450 1.0000000 0.1449678 0.1371596 0.1399074
## [3,] 0.1457634 0.1449678 1.0000000 0.2157428 0.1530311
## [4,] 0.1514811 0.1371596 0.2157428 1.0000000 0.1459009
## [5,] 0.2030989 0.1399074 0.1530311 0.1459009 1.0000000
```

4) Simulating paths from the full model

Here we will start from the fitted marginal distributions and t copula to simulate new random vectors of size m and dimension 5. We will repeat this simulation $B = 200$ times to get an idea of its variability. The simulation steps are explicated in the R code below.

```

B <- 200
m <- ceiling(n/10) # length of the simulated paths
X.lst <- lapply(1:B, function(b) {
  ## 1) Simulate from the fitted copula
  U. <- rCopula(m, copula = tc)

  ## 2) Quantile-transform to standardized t distributions (for ugarchsim())
  Z. <- sapply(1:p, function(j) sqrt((nu.mar[j]-2)/nu.mar[j]) * qt(U.[,j], df = nu.mar[j]))

  ## 3) Use these multivariate dependent t innovations to sample from the time series
  sim <- lapply(1:p, function(j)
    ugarchsim(fits[[j]], n.sim = m, m.sim = 1,
      custom.dist = list(name = "sample",
        distfit = Z.[,j, drop = FALSE])))

  ## 4) Extract simulated series
  sapply(sim, function(x) fitted(x)) # simulated multivariate series X_t x@simulation$seriesSim
})

```

Which generates a List of length B containing (n x p)-matrices

5) Predict the aggregated loss and $\text{VaR}_{0.99}$

5.1) Compute the aggregated loss by summing up the losses of all $p = 5$ constituents at each time t in the original data X .

5.2) Compute the aggregated losses for each simulated matrix in the list 'X.lst': We thus obtain B instances of the simulated losses over time.

The predicted loss is the average over all simulated losses. Compute an α -confidence interval by taking the $\alpha/2$ and $1 - \alpha/2$ empirical quantiles of the simulated losses.

5.3) Compute the $\text{VaR}_{0.99}$ based on the simulated losses, and plot them together with the original return data.

```

# aggregated loss; n-vector
Xs <- rowSums(X)

# simulated aggregated losses; (m, B)-matrix
Xs. <- sapply(X.lst, rowSums)

# predicted aggregated loss; m-vector
Xs.mean <- rowMeans(Xs.)

# CIs; (2, m)-matrix
Xs.CI <- apply(Xs., 1, function(x) quantile(x, probs = c(0.025, 0.975)))

# confidence level
alpha <- 0.99

# VaR_alpha; m-vector
VaR <- apply(Xs., 1, function(x) quantile(x, probs = alpha))


tm <- index(SP500_const)
start <- match(time[1], as.character(tm))
past <- tm[start:(start+n-1)]

```

```

future <- tm[(start+n):(start+n+m-1)]
plot(past, Xs, type = "l", xlim = range(c(past, future)), xlab = "", ylab = "") # actual (past) losses
polygon(c(future, rev(future)), c(Xs.CI[1,], rev(Xs.CI[2,])),
        border = NA, col = "grey80") # CI region
lines(future, Xs.mean, col = "royalblue3") # predicted aggregated loss
lines(future, Xs.CI[1,], col = "grey50") # lower CI
lines(future, Xs.CI[2,], col = "grey50") # upper CI
lines(future, VaR, col = "maroon3") # VaR_alpha
legend("bottomright", bty = "n", lty = rep(1, 4),
       col = c("black", "royalblue3", "grey50", "maroon3"),
       legend = c("(Aggregated) loss", "(Simulated) predicted loss",
                  "95% CIs", as.expression(substitute("Simulated"~VaR[a], list(a = alpha))))))

```



References This tutorial is based on R code from QRMtutorials.org